

Copy_of_Keras_Mnist

March 4, 2020

0.1 Keras -- MLPs on MNIST

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
        from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
```

```
In [0]: %matplotlib notebook
        import matplotlib.pyplot as plt
        import numpy as np
        import time
        # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
        # https://stackoverflow.com/a/14434334
        # this function is used to update the plots for each epoch and error
        def plt_dynamic(x, vy, ty, ax, colors=['b']):
            ax.plot(x, vy, 'b', label="Validation Loss")
            ax.plot(x, ty, 'r', label="Train Loss")
            plt.legend()
            plt.grid()
            fig.canvas.draw()
```

```
In [0]: %matplotlib inline
```

```
In [0]: # the data, shuffled and split between train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [6]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)
        print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
        # for each image we have a (28*28) vector
        # we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [8]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)")
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
In [9]: # An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

```

55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 16
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

```

```

In [0]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data
        #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

```

```

X_train = X_train/255
X_test = X_test/255

```

```

In [11]: # example data point after normlizing
         print(X_train[0])

```

```

[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.1176471 0.07058824 0.07058824 0.07058824
 0.49411765 0.53333333 0.68627451 0.10196078 0.65098039 1.
 0.96862745 0.49803922 0.      0.      0.      0.]

```

0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

5

[illegible]

```
In [12]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

```
In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/
```

```

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_r
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [0]: # some model parameters

```

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [0]: # start building a model

```

model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

```

```

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

In [17]: # Before training a model, you need to configure the learning process, which is done by

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector with a 1
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_data=None,
# shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset)

# it returns A History object. Its History.history attribute is a record of training metrics values at successive epochs, as well as validation loss values and validation metrics values at successive epochs

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:144: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/math_ops.py:396: to_float32 is deprecated and will be removed in a future version. Instructions for updating:
Use tf.float32 in 2.x
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:144: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:144: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.

```


WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:397: *tf.nn.conv2d* is deprecated and will be removed in a future version.

60000/60000 [=====] - 11s 180us/step - loss: 1.2521 - acc: 0.7184 - val_loss: 1.2521

Epoch 2/20

60000/60000 [=====] - 1s 24us/step - loss: 0.7087 - acc: 0.8454 - val_loss: 0.7087

Epoch 3/20

60000/60000 [=====] - 1s 24us/step - loss: 0.5826 - acc: 0.8623 - val_loss: 0.5826

Epoch 4/20

60000/60000 [=====] - 1s 24us/step - loss: 0.5221 - acc: 0.8703 - val_loss: 0.5221

Epoch 5/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4852 - acc: 0.8758 - val_loss: 0.4852

Epoch 6/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4598 - acc: 0.8803 - val_loss: 0.4598

Epoch 7/20

60000/60000 [=====] - 1s 25us/step - loss: 0.4410 - acc: 0.8842 - val_loss: 0.4410

Epoch 8/20

60000/60000 [=====] - 1s 25us/step - loss: 0.4263 - acc: 0.8869 - val_loss: 0.4263

Epoch 9/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4144 - acc: 0.8892 - val_loss: 0.4144

Epoch 10/20

60000/60000 [=====] - 1s 24us/step - loss: 0.4046 - acc: 0.8916 - val_loss: 0.4046

Epoch 11/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3962 - acc: 0.8931 - val_loss: 0.3962

Epoch 12/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3891 - acc: 0.8946 - val_loss: 0.3891

Epoch 13/20

60000/60000 [=====] - 1s 25us/step - loss: 0.3828 - acc: 0.8958 - val_loss: 0.3828

Epoch 14/20

60000/60000 [=====] - 1s 23us/step - loss: 0.3772 - acc: 0.8972 - val_loss: 0.3772

Epoch 15/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3722 - acc: 0.8978 - val_loss: 0.3722

Epoch 16/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3677 - acc: 0.8991 - val_loss: 0.3677

Epoch 17/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3637 - acc: 0.8997 - val_loss: 0.3637

Epoch 18/20

```

60000/60000 [=====] - 1s 24us/step - loss: 0.3599 - acc: 0.9009 - val.
Epoch 19/20
60000/60000 [=====] - 1s 25us/step - loss: 0.3566 - acc: 0.9017 - val.
Epoch 20/20
60000/60000 [=====] - 1s 25us/step - loss: 0.3534 - acc: 0.9025 - val.

```

```

In [74]: score = model.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

         # we will get val_loss and val_acc only when you pass the paramter validation_data
         # val_loss : validation loss
         # val_acc : validation accuracy


         # loss : training loss
         # acc : train accuracy
         # for each key in history.histrory we will have a list of length equal to number of

         vy = history.history['val_loss']
         ty = history.history['loss']
         plt_dynamic(x, vy, ty, ax)

```

Test score: 0.335641682690382

Test accuracy: 0.9094



MLP + ReLu activation + Adam Optimizer

In [31]: # *Multilayer perceptron*

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(226, activation='relu', input_shape=(input_dim,)))
model_sigmoid.add(Dense(132, activation='relu'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 226)	177410
dense_14 (Dense)	(None, 132)	29964
dense_15 (Dense)	(None, 10)	1330
Total params: 208,704		
Trainable params: 208,704		
Non-trainable params: 0		

```

In [32]: model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc'])

        history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 38us/step - loss: 0.2741 - acc: 0.9215 - val_
Epoch 2/20
60000/60000 [=====] - 2s 32us/step - loss: 0.1078 - acc: 0.9671 - val_
Epoch 3/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0704 - acc: 0.9785 - val_
Epoch 4/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0525 - acc: 0.9838 - val_
Epoch 5/20
60000/60000 [=====] - 2s 30us/step - loss: 0.0383 - acc: 0.9881 - val_
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0303 - acc: 0.9904 - val_
Epoch 7/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0231 - acc: 0.9925 - val_
Epoch 8/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0189 - acc: 0.9941 - val_
Epoch 9/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0163 - acc: 0.9944 - val_
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0138 - acc: 0.9955 - val_
Epoch 11/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0157 - acc: 0.9947 - val_
Epoch 12/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0113 - acc: 0.9965 - val_
Epoch 13/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0087 - acc: 0.9971 - val_
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0108 - acc: 0.9963 - val_
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0104 - acc: 0.9962 - val_
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0060 - acc: 0.9979 - val_
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0083 - acc: 0.9971 - val_
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0072 - acc: 0.9976 - val_
Epoch 19/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0071 - acc: 0.9974 - val_
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0094 - acc: 0.9968 - val_

In [73]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])

```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.histrory we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09268326033384333

Test accuracy: 0.9798



```

In [72]: w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 2 layers

In [35]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
model_drop = Sequential()

model_drop.add(Dense(226, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(132, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 226)	177410
batch_normalization_3 (Batch Normalization)	(None, 226)	904
dropout_3 (Dropout)	(None, 226)	0
dense_17 (Dense)	(None, 132)	29964
batch_normalization_4 (Batch Normalization)	(None, 132)	528
dropout_4 (Dropout)	(None, 132)	0
dense_18 (Dense)	(None, 10)	1330
Total params: 210,136		
Trainable params: 209,420		
Non-trainable params: 716		

In [71]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ver
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 126us/step - loss: 0.2252 - acc: 0.9416 - val.
Epoch 2/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2182 - acc: 0.9431 - val.
Epoch 3/20
60000/60000 [=====] - 6s 93us/step - loss: 0.2109 - acc: 0.9455 - val.
Epoch 4/20
60000/60000 [=====] - 6s 95us/step - loss: 0.2110 - acc: 0.9453 - val.
Epoch 5/20
60000/60000 [=====] - 6s 93us/step - loss: 0.2046 - acc: 0.9467 - val.
Epoch 6/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1994 - acc: 0.9486 - val.
Epoch 7/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1941 - acc: 0.9502 - val.
Epoch 8/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1912 - acc: 0.9513 - val.
Epoch 9/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1901 - acc: 0.9510 - val.
Epoch 10/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1838 - acc: 0.9525 - val.
Epoch 11/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1849 - acc: 0.9516 - val.
Epoch 12/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1808 - acc: 0.9528 - val.
Epoch 13/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1806 - acc: 0.9520 - val.
Epoch 14/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1726 - acc: 0.9549 - val.
Epoch 15/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1772 - acc: 0.9540 - val.
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1736 - acc: 0.9540 - val.
Epoch 17/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1700 - acc: 0.9561 - val.
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1668 - acc: 0.9564 - val.
Epoch 19/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1650 - acc: 0.9571 - val.
Epoch 20/20
60000/60000 [=====] - 5s 92us/step - loss: 0.1665 - acc: 0.9567 - val.
```

```
In [70]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
```



```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.histrory we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12922746661901474

Test accuracy: 0.9667



```

In [69]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 3 layers

In [44]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(442, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(123, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(82, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_23 (Dense)	(None, 442)	346970
batch_normalization_8 (Batch Normalization)	(None, 442)	1768
dropout_8 (Dropout)	(None, 442)	0
dense_24 (Dense)	(None, 123)	54489
batch_normalization_9 (Batch Normalization)	(None, 123)	492
dropout_9 (Dropout)	(None, 123)	0
dense_25 (Dense)	(None, 82)	10168
batch_normalization_10 (Batch Normalization)	(None, 82)	328
dropout_10 (Dropout)	(None, 82)	0

```

-----
dense_26 (Dense)                (None, 10)                830
=====

```

```

Total params: 415,045
Trainable params: 413,751
Non-trainable params: 1,294
-----

```

```

In [45]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```

        history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))

```

```

Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/20
60000/60000 [=====] - 5s 87us/step - loss: 0.8982 - acc: 0.7228 - val_loss: 0.7100
Epoch 2/20
60000/60000 [=====] - 4s 67us/step - loss: 0.4125 - acc: 0.8783 - val_loss: 0.3800
Epoch 3/20
60000/60000 [=====] - 4s 67us/step - loss: 0.3221 - acc: 0.9079 - val_loss: 0.3200
Epoch 4/20
60000/60000 [=====] - 4s 67us/step - loss: 0.2656 - acc: 0.9241 - val_loss: 0.2800
Epoch 5/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2327 - acc: 0.9330 - val_loss: 0.2500
Epoch 6/20
60000/60000 [=====] - 4s 67us/step - loss: 0.2128 - acc: 0.9402 - val_loss: 0.2300
Epoch 7/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1942 - acc: 0.9448 - val_loss: 0.2100
Epoch 8/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1787 - acc: 0.9493 - val_loss: 0.1900
Epoch 9/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1686 - acc: 0.9521 - val_loss: 0.1700
Epoch 10/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1577 - acc: 0.9548 - val_loss: 0.1500
Epoch 11/20
60000/60000 [=====] - 4s 70us/step - loss: 0.1524 - acc: 0.9563 - val_loss: 0.1400
Epoch 12/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1422 - acc: 0.9597 - val_loss: 0.1300
Epoch 13/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1357 - acc: 0.9620 - val_loss: 0.1200
Epoch 14/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1305 - acc: 0.9627 - val_loss: 0.1100
Epoch 15/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1242 - acc: 0.9642 - val_loss: 0.1000
Epoch 16/20
60000/60000 [=====] - 4s 69us/step - loss: 0.1178 - acc: 0.9659 - val_loss: 0.0900
Epoch 17/20
60000/60000 [=====] - 4s 65us/step - loss: 0.1110 - acc: 0.9682 - val_loss: 0.0800

```

```

Epoch 18/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1122 - acc: 0.9675 - val.
Epoch 19/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1022 - acc: 0.9705 - val.
Epoch 20/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1047 - acc: 0.9708 - val.

```

```

In [66]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_test, Y_test))

         # we will get val_loss and val_acc only when you pass the paramter validation_data
         # val_loss : validation loss
         # val_acc : validation accuracy


         # loss : training loss
         # acc : train accuracy
         # for each key in history.history we will have a list of length equal to number of epochs

         vy = history.history['val_loss']
         ty = history.history['loss']
         plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12922746661901474

Test accuracy: 0.9667



```
In [65]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)
```


```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='m')
plt.xlabel('Output Layer')
```

```
plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 5 layers

In [48]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

```

```

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

```

```

model_drop.add(Dense(output_dim, activation='softmax'))

```

```

model_drop.summary()

```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 128)	100480
batch_normalization_11 (Batch Normalization)	(None, 128)	512
dropout_11 (Dropout)	(None, 128)	0
dense_28 (Dense)	(None, 128)	16512
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dropout_12 (Dropout)	(None, 128)	0
dense_29 (Dense)	(None, 128)	16512
batch_normalization_13 (Batch Normalization)	(None, 128)	512
dropout_13 (Dropout)	(None, 128)	0
dense_30 (Dense)	(None, 128)	16512
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_14 (Dropout)	(None, 128)	0
dense_31 (Dense)	(None, 128)	16512
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dropout_15 (Dropout)	(None, 128)	0


```

-----
dense_32 (Dense)                (None, 10)                1290
=====
Total params: 170,378
Trainable params: 169,098
Non-trainable params: 1,280
-----

```

```
In [49]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))
```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 7s 124us/step - loss: 2.0302 - acc: 0.3242 - val_loss: 1.8750
Epoch 2/20
60000/60000 [=====] - 6s 94us/step - loss: 1.0387 - acc: 0.6370 - val_loss: 0.8750
Epoch 3/20
60000/60000 [=====] - 5s 91us/step - loss: 0.7431 - acc: 0.7597 - val_loss: 0.6250
Epoch 4/20
60000/60000 [=====] - 6s 92us/step - loss: 0.6024 - acc: 0.8151 - val_loss: 0.5000
Epoch 5/20
60000/60000 [=====] - 5s 91us/step - loss: 0.5116 - acc: 0.8481 - val_loss: 0.4375
Epoch 6/20
60000/60000 [=====] - 6s 92us/step - loss: 0.4611 - acc: 0.8689 - val_loss: 0.3750
Epoch 7/20
60000/60000 [=====] - 6s 93us/step - loss: 0.4127 - acc: 0.8841 - val_loss: 0.3125
Epoch 8/20
60000/60000 [=====] - 5s 92us/step - loss: 0.3831 - acc: 0.8939 - val_loss: 0.2500
Epoch 9/20
60000/60000 [=====] - 6s 98us/step - loss: 0.3468 - acc: 0.9057 - val_loss: 0.1875
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.3286 - acc: 0.9119 - val_loss: 0.1250
Epoch 11/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3094 - acc: 0.9169 - val_loss: 0.0625
Epoch 12/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2925 - acc: 0.9216 - val_loss: 0.0000
Epoch 13/20
60000/60000 [=====] - 6s 94us/step - loss: 0.2782 - acc: 0.9266 - val_loss: 0.0000
Epoch 14/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2696 - acc: 0.9290 - val_loss: 0.0000
Epoch 15/20
60000/60000 [=====] - 6s 93us/step - loss: 0.2603 - acc: 0.9310 - val_loss: 0.0000
Epoch 16/20
60000/60000 [=====] - 5s 90us/step - loss: 0.2522 - acc: 0.9344 - val_loss: 0.0000
Epoch 17/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2510 - acc: 0.9347 - val_loss: 0.0000

```

```
Epoch 18/20
60000/60000 [=====] - 5s 91us/step - loss: 0.2356 - acc: 0.9386 - val.
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.2358 - acc: 0.9390 - val.
Epoch 20/20
60000/60000 [=====] - 6s 99us/step - loss: 0.2286 - acc: 0.9405 - val.
```

```
In [64]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1,nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_test, Y_test))


        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of epochs

        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12922746661901474

Test accuracy: 0.9667



```
In [63]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
```

```


plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='purple')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='orange')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Summary

```

In [0]: from prettytable import PrettyTable
        summary = PrettyTable()

In [0]: summary.field_names = ["Model", "Dropout", "Test Loss", "Test Accuracy"]

In [77]: summary.add_row(["MLP + ReLu activation + Adam Optimizer", "0.5", "0.092", "0.9798"])
        summary.add_row(["MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 2

```

```
summary.add_row(["MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 3 layers"])
summary.add_row(["MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 5 layers"])

print(summary)
```

Model	Dropout	Test Loss
MLP + ReLu activation + Adam Optimizer	0.5	0.09
MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 2 layers	0.5	0.12
MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 3 layers	0.5	0.12
MLP + ReLu activation + BatchNormalisation + Dropout + ADAM with 5 layers	0.5	0.12